Event-Time Geometry: A Critical Analysis of its Foundations, Novelty, and Formalization for Asynchronous Systems

Foundational Principles of Event-Time Geometry

The analysis of modern asynchronous and event-driven computational systems, from distributed networks to neuromorphic hardware, demands a formal framework capable of reasoning about causality with metric precision. The Event-Time Geometry (ETG) framework offers such a formalism by drawing a direct and productive analogy to the geometry of spacetime in physics. To ground the subsequent analysis of its novelty, limitations, and potential for formalization, it is essential to first establish a self-contained exposition of its core principles, beginning with its idealized deterministic formulation and progressing to its more realistic probabilistic extension.

The Deterministic Framework (D-ETG): An Analogy to Spacetime Physics

Deterministic Event-Time Geometry (D-ETG) provides a foundational, idealized model of causality under conditions of perfect predictability and knowledge. It establishes a world of discrete, point-like occurrences whose positions in both space and time are known with absolute certainty, forming a clear baseline against which the complexities of real-world uncertainty can be measured.

Core Components

The D-ETG framework is constructed upon a manifold composed of several fundamental elements. The first is the set of **Locations**, denoted by L. A location $l \in L$ is a spatial identifier, which could represent a processor ID in a distributed system, a neuron index in a neuromorphic network, or a sensor pixel in an event-based camera. This set is endowed with a distance function, $d: L \times L \rightarrow \mathbb{R}_{\ge 0}$, which satisfies the standard axioms of a metric and quantifies the spatial separation or, more abstractly, the minimum communication latency between any two locations.

The second element is the set of **Timestamps**, T, which constitutes the time domain. This is typically represented by the set of real numbers, $T = \mathbb{R}$, ensuring that every event can be assigned a precise, real-valued time coordinate.

Upon this space-time manifold, the fundamental ontological unit is the **Event**. An event e is conceptualized as a perfectly localized point, defined by a tuple $e = (l, t)$ where $l \in L$ and $t \in T$. Some definitions also include a payload or effect, ϕ, making an event a tuple $e = (l, t, \phi)$, though the geometric properties primarily concern the location and time coordinates.

The Geometry of Causality

D-ETG imposes a strict partial order on the set of events that is governed not by pure logic, but by the physical or systemic constraints of information propagation. An event $e_1 = (l_1, t_1)$ can causally influence a subsequent event $e_2 = (l_2, t_2)$, a relationship denoted as $e_1 \rightarrow e_2$, only if there is sufficient time for a signal to travel from location l_1 to l_2. This is formalized by a fundamental causal delay constraint. If a maximum signal propagation speed within the system is defined as c, then a necessary condition for $e_1 \rightarrow e_2$ is given by the inequality :
This inequality asserts that the temporal separation, $\Delta t = t_2 - t_1$, must be at least as great as the minimum time required to traverse the spatial distance, $\Delta x = d(l_1, l_2)$, at the maximum speed c. If this condition is violated, e_1 is causally prohibited from influencing e_2. This geometric view of causality represents a significant conceptual advance over purely logical or topological frameworks for event ordering, such as Lamport's "happened-before" relation. Lamport's model defines a partial order based on the sequence of events within a single process and the sending and receiving of messages, allowing it to determine if one event logically preceded another. However, it lacks any notion of metric distance or real time; it cannot answer questions like, "By how much time did A precede B?" or "Could A still have caused B if the network latency were higher?". D-ETG, by introducing the metric components $d(l_1, l_2)$ and c, grounds the causal relationship in the physical reality of the system, enabling quantitative analysis of properties like latency, temporal slack, and performance budgets—prerequisites for the design of systems with hard real-time constraints.

The Minkowski-like Interval and Causal Cones

To quantify the separation between events in a manner that respects the underlying causal structure, D-ETG introduces an invariant measure analogous to the spacetime interval in Minkowski space. Given two events, $e_1 = (l_1, t_1)$ and $e_2 = (l_2, t_2)$, with temporal separation $\Delta t = t_2 - t_1$ and spatial separation $\Delta x = d(l_1, l_2)$, the squared event-time interval is defined as :
This single scalar quantity provides a powerful, tripartite classification of the relationship between any pair of events, directly reflecting their causal connectability :

- **Timelike Separation** ($(\Delta s)^2 > 0$): When the temporal separation is large enough to overcome the spatial distance, the events are causally connectable. Assuming $\Delta t > 0$, event e_2 lies within the future causal cone of e_1, meaning e_1 could have influenced e_2.
- **Spacelike Separation** ($(\Delta s)^2 < 0$): When the spatial distance is too great for the given temporal separation, the events are causally disconnected. No signal traveling at or below speed c could have connected them. They exist outside of each other's causal cones.
- **Lightlike Separation** ($(\Delta s)^2 = 0$): This boundary case represents events connected by a signal traveling at precisely the maximum speed c. Event e_2 lies on the edge of e_1's future causal cone.

This classification is absolute and exhaustive, partitioning all event pairs based on their intrinsic causal relationship. The **causal cone** of an event e is thus defined as the set of all points in the event-time manifold that have a timelike or lightlike separation from e, cleanly delineating the region of potential causal influence from the region of causal independence. These cones have sharp, well-defined boundaries, mirroring the light cones of relativistic physics.

The Event-Time Invariance Theorem

The theoretical cornerstone of D-ETG is the Event-Time Invariance Theorem, which states that the value of the squared event-time interval, $(\Delta s)^2$, is invariant under a specific class of coordinate transformations known as Lorentz-like transformations. In the context of distributed systems, a "frame of reference" can be thought of as a particular clock synchronization scheme or the perspective of a specific observer or subsystem. A transformation between two frames F and F' moving at a relative "velocity" v can be modeled as :
where $\gamma = 1 / \sqrt{1 - v^2/c^2}$.
The theorem guarantees that while different observers (i.e., different frames) may disagree on the individual time and space coordinates of events—a phenomenon analogous to the relativity of simultaneity in physics—they will all calculate the exact same value for the interval $(\Delta s)^2$ between any two events. The profound implication of this invariance is that the causal structure of the system is absolute. All observers, regardless of their reference frame, will agree on the classification of an event pair as timelike, spacelike, or lightlike. The question of whether two events can be causally related is an objective, frame-independent truth.

The Probabilistic Extension (P-ETG): Embracing Uncertainty

The deterministic nature of D-ETG is simultaneously its greatest strength—providing a powerful and consistent framework for analyzing predictable systems—and its most significant weakness in practical application. Its core assumptions of perfectly accurate clocks, zero transmission jitter, and constant, known propagation delays are systematically violated in the asynchronous, noisy, and complex environments of modern computational systems. This inherent "brittleness" motivates and necessitates the development of a probabilistic extension, Probabilistic Event-Time Geometry (P-ETG). P-ETG achieves this by systematically replacing each deterministic concept with a stochastic counterpart, moving from a world of precise points and sharp boundaries to one of probability distributions and graded confidence.

From Points to Clouds

The most fundamental shift in P-ETG is the redefinition of the event itself. An event is no longer a deterministic point (l, t) but is instead represented as a random variable, $E = (L_E, T_E)$, characterized by a joint probability distribution, $P_E(l, t)$, over the space-time manifold. This distribution can be visualized as a "probability cloud" or "blob" in a space-time diagram, where the density of the cloud at any point (l, t) corresponds to the likelihood that the event occurred at that specific location and time.
This representation directly captures real-world uncertainty. For example, the timestamp of a neuron's spike might be modeled as a Gaussian random variable, $T_E \sim \mathcal{N}(\mu_t, \sigma_t^2)$, where μ_t is the expected firing time and σ_t^2 is the variance due to intrinsic neural noise or clock jitter. This "event cloud" picture acknowledges that our knowledge of an event's coordinates is often limited to a set of statistical properties or bounded intervals. In the limit as the variance approaches zero, the probabilistic event converges to its deterministic, point-like counterpart.

Stochastic Intervals and Fuzzy Causality

The introduction of uncertainty at the level of event coordinates has profound consequences for the entire geometric structure. Since the event coordinates (L_1, T_1) and (L_2, T_2) are now random variables, the temporal separation $\Delta T = T_2 - T_1$ and the spatial separation $\Delta X = d(L_1, L_2)$ also become random variables. Consequently, the event-time interval, now a stochastic quantity $(\Delta S)^2 = c^2(\Delta T)^2 - (\Delta X)^2$, is itself a random variable with a full probability distribution over possible interval values, $P((\Delta S)^2)$.
This fundamentally changes the nature of causal classification. The sign of $(\Delta S)^2$, which deterministically partitions event pairs, is no longer certain. An event pair might have a 90% probability of being timelike separated and a 10% probability of being spacelike separated. This leads to the concept of **fuzzy causal classes**, where the relationship between events is described by a set of probabilities rather than a single, crisp label :

-
-
-

This probabilistic classification provides a much more nuanced and realistic depiction of causality in systems where timing is not perfectly precise.

The Probabilistic Causal Relation

With the deterministic causal condition rendered uncertain, the binary true/false nature of the causal relation $e_1 \rightarrow e_2$ is replaced by a graded measure of confidence. The causal influence probability is defined as the probability that the necessary condition for causality holds, given the distributions of the event times and locations :
This value, ranging from 0 to 1, represents our degree of belief that e_1 could have causally affected e_2. A probability near 1 indicates high confidence in a potential causal link, while a probability near 0 suggests that e_2 almost certainly occurred too soon for e_1 to have been a cause. This formalism has been explored in recent research through relations like "likely-happened-before," which explicitly assigns a probability to the happened-before relation, extending Lamport's framework to order events that would be considered concurrent in a deterministic view. The trade-off is that this probabilistic relation may not be transitive; it is possible that $P(A \rightarrow B)$ and $P(B \rightarrow C)$ are both high, while $P(A \rightarrow C)$ is not. As a result, the causal structure of the system is no longer a strict partial order but is more accurately represented as a weighted directed acyclic graph.

Fuzzy Lightcones and Stochastic Invariance

The sharp, well-defined boundary of the deterministic causal cone dissolves under uncertainty into a "fuzzy" or blurred transition region. This can be visualized as a cone whose edges are smeared out, reflecting the probabilistic nature of causality at the boundary. An event that would be classically just outside the lightcone might now have a small but non-zero probability of being influenced, due to timing fluctuations. This fuzzy lightcone can be formalized by defining probabilistic causal cones based on confidence levels. For example, the 90% causal cone of an event e_1 could be defined as the set of all space-time points (l, t) for which the causal influence probability $P(e_1 \rightarrow (l,t))$ is greater than or equal to 0.9.
Just as invariance is the bedrock of D-ETG, a corresponding principle of **stochastic invariance** is crucial for P-ETG. Since the interval $(\Delta S)^2$ is now a random variable, its specific value

cannot be preserved across reference frames. Instead, stochastic Lorentz invariance demands that the *statistical law*, or the full probability distribution, of the interval is invariant. This ensures that any physically meaningful probabilistic statement, such as the value of $P(e_1 \rightarrow e_2)$, is frame-independent. If one observer calculates a 95% chance of a causal link, all other observers in different inertial frames must arrive at the same conclusion.
This transition from a deterministic to a probabilistic framework represents a fundamental shift from a "God's-eye view" of a system, where all state variables are perfectly known, to an observer-centric perspective, where knowledge is fundamentally incomplete and must be described by probability distributions. The "fuzzy lightcone" is conceptually analogous to quantum phenomena like tunneling, where a particle has a non-zero probability of existing in a classically forbidden state. This reframes P-ETG not merely as an engineering tool for handling noise, but as a deeper, more physically realistic model of information and causality.
Consequently, the notion of "correctness" in system design is transformed. In D-ETG, correctness is absolute: a deadline is either met or missed. In P-ETG, correctness becomes a probabilistic guarantee: for example, "the probability of a deadline miss is less than 10^{-9}". This paradigm shift compels system designers to engage in explicit risk management, quantitatively balancing performance gains against the calculable probability of failure, a concept central to modern fields like probabilistic scheduling.

Novelty Analysis and Positioning within the Theoretical Landscape

Probabilistic Event-Time Geometry does not exist in isolation; it is part of a broader landscape of formalisms designed to reason about time and causality in computational systems. To fully appreciate its unique contributions and establish its novelty, it is essential to position P-ETG relative to other well-known frameworks. This comparative analysis reveals a spectrum of models, each with distinct trade-offs in terms of expressive power, computational cost, and the nature of the guarantees they provide. P-ETG emerges at the most expressive end of this spectrum, offering the richest model of uncertainty at the cost of greater complexity.

Beyond Logical and Bounded Time in Distributed Systems

The foundational work in ordering events in distributed systems provides a crucial point of contrast, highlighting ETG's introduction of metric and probabilistic reasoning.

Versus Lamport and Vector Clocks

The pioneering work of Leslie Lamport introduced the "happened-before" ($\rightarrow$) relation, a purely logical framework for capturing potential causality in asynchronous distributed systems. A Lamport clock, a simple counter maintained by each process, ensures that if event A happened-before event B, then the logical timestamp of A is less than that of B. Vector clocks extend this to provide an "if and only if" condition, creating an isomorphism between the causal partial order and the order of the vector timestamps.
However, the nature of these logical clocks is topological, not metric. They establish a partial order but contain no information about real-time duration or physical separation. Events for which no causal path exists are deemed "concurrent," with no way to distinguish between events that occurred nanoseconds apart and those that occurred hours apart. P-ETG provides a

significant extension in two key ways:

1. **Metric Foundation:** By incorporating a physical distance $d(l_1, l_2)$ and a maximum signal speed c, P-ETG introduces a metric that quantifies the "gap" between events. The event-time interval $(\Delta s)^2$ is a real-valued measure of separation, not just a logical ordering.
2. **Probabilistic Ordering of Concurrent Events:** P-ETG directly addresses the ambiguity of concurrency. For two events that are concurrent under Lamport's definition, P-ETG can leverage knowledge of clock offset distributions to compute the probability that one truly occurred before the other. The "likely-happened-before" relation, $A \xrightarrow{p} B$, states that event A happened before event B with probability p. This allows the system to establish a total order with a specified degree of confidence, resolving ambiguity where logical clocks can only report concurrency.

Versus Google's TrueTime

Google's Spanner database introduced TrueTime, a groundbreaking approach to managing time uncertainty in a globally distributed system. TrueTime does not provide a single timestamp but rather a bounded interval, $[t_{\text{earliest}}, t_{\text{latest}}]$, which is guaranteed to contain the true, absolute time of an event. This interval-based representation allows for the implementation of externally consistent transactions; if the interval for transaction T_1 ends before the interval for transaction T_2 begins, their order is certain and globally agreed upon.

TrueTime is a powerful model for building systems that require strong correctness guarantees by reasoning about worst-case uncertainty. However, its approach is inherently conservative. If the uncertainty intervals of two events overlap, TrueTime can make no definitive statement about their order; they must be treated as potentially concurrent, often forcing a "commit wait" to ensure correctness.

P-ETG offers a more granular and potentially more performant alternative by modeling uncertainty with a full probability distribution rather than a hard interval. While TrueTime's intervals might overlap, P-ETG can still compute the probability $P(t_A < t_B)$ by convolving the respective time distributions of events A and B. For example, even if their 99.99% confidence intervals overlap, the probability that A occurred before B might still be extremely high (e.g., 99.9999%). A system based on P-ETG can leverage this information to make more aggressive optimizations. In applications like fair scheduling, it can choose to order A before B, accepting a quantifiable and potentially infinitesimal risk of misordering in exchange for higher throughput or lower latency. TrueTime, by contrast, would be forced into a more conservative and potentially slower execution path to avoid any possibility of misordering.

Distinctions from Causal Set Theory (CST)

While both ETG and Causal Set Theory (CST) draw inspiration from the causal structure of spacetime, they represent fundamentally different approaches to modeling events and causality, operating at different levels of abstraction and with different primary goals.

Continuous Manifold vs. Discrete Structure

The most profound distinction lies in the assumed structure of the underlying space. ETG, in both its deterministic and probabilistic forms, presupposes a continuous underlying space-time manifold, analogous to the continuum of classical relativity. Events are points or distributions *on*

this pre-existing manifold.
In stark contrast, Causal Set Theory posits that spacetime is fundamentally discrete at the most microscopic level (the Planck scale). A causal set (or "causet") is a locally finite, partially ordered set $(C, \prec)$ where the elements of C are the discrete "atoms" of spacetime, and the order relation $\prec$ represents the causal relationship between them. In CST, there is no underlying continuous manifold; geometry itself is an emergent property derived from the combinatorial and statistical properties of the causet, a principle Sorkin coined as "Order + Number = Geometry". While a causet can be generated by a random "sprinkling" of points onto a continuous manifold for modeling purposes, the resulting causet retains no geometric information beyond the order relation.

Primacy of Causality vs. Metric

This structural difference leads to a different logical priority. In ETG, the Minkowski-like metric is the primary object. The distance function d and the constant c are foundational, and the causal structure—the classification of event pairs as timelike, spacelike, or lightlike—is *derived* from the sign of the interval $(\Delta s)^2$.
In CST, the causal order relation $\prec$ is the sole fundamental structure. Concepts analogous to distance and time are *derived* from this order. For instance, a "timelike geodesic" between two elements in a causet is analogized by the longest chain of relations connecting them, and the "proper time" is taken to be proportional to the number of elements in this chain. The metric is emergent, not axiomatic.

Nature of Randomness

The role of probability also differs significantly. In P-ETG, the background manifold is fixed and deterministic. Randomness is introduced in the *coordinates* of events on this background, representing measurement uncertainty, jitter, or other forms of noise in a macroscopic system.
In CST, randomness is typically introduced through the process of Poisson sprinkling, which is a procedure for generating a causet that faithfully approximates a given continuous spacetime. The randomness is in the selection of the discrete elements that constitute the causet itself. The goal of CST as a theory of quantum gravity is to define a path integral or "sum over histories" by summing over all possible causal sets, a fundamentally different endeavor from modeling stochasticity in a classical computational system.
This comparison positions ETG as a "mesoscopic" model. It is more physically realistic than purely abstract logical models like Lamport clocks but avoids the deep, unresolved questions of quantum gravity and Planck-scale physics inherent in a fundamental theory like CST. It pragmatically adopts the mathematical structure of classical relativity—a continuous manifold with a metric—and applies it as a rigorous, powerful metaphor to the domain of computation, where c is not the speed of light but a system-specific maximum information propagation speed. This allows system designers to use the well-developed intuitions of spacetime physics to reason about computational causality without needing to solve quantum gravity.

Relationship to Spatio-Temporal Point Processes (STPPs)

Temporal point processes (TPPs), and their spatio-temporal extensions (STPPs), are a class of stochastic models used to describe the occurrence of discrete events in continuous time and space. A key class of such models are Hawkes processes, which are characterized by their self-exciting nature: the occurrence of an event temporarily increases the rate, or conditional intensity $\lambda(t, s)$, of future events. This makes them exceptionally well-suited for modeling phenomena with temporal clustering, such as aftershocks following an earthquake or bursts of neural firing.

ETG as a Structural Prior

In standard STPP models, the influence of a past event on the future intensity is governed by a kernel function, $\phi(\Delta t, \Delta s)$, which typically decays with temporal and spatial distance. However, the choice of this kernel is often a matter of statistical convenience (e.g., an exponential or power-law decay) and does not inherently respect the physical constraints of causality. A standard Hawkes process could be parameterized to model an event in one location triggering another faster than the system's information propagation speed would allow; the model would simply learn a near-zero influence coefficient for such physically impossible interactions.

P-ETG offers a more principled and physically motivated constraint on this influence. The causal cone structure inherited from D-ETG imposes a hard, non-negotiable constraint on the influence kernel of any point process model built upon it: the influence of an event at (l_1, t_1) on the intensity at any point (l_2, t_2) is strictly zero if the two points are spacelike separated. This can be expressed as:

If $c^2(t_2 - t_1)^2 - d(l_1, l_2)^2 < 0$, then the contribution of the event at (l_1, t_1) to $\lambda(t_2, l_2)$ is zero.

This integration of a hard causal geometry provides a powerful **structural prior** that can improve model accuracy, reduce the number of parameters to be learned, and ensure physical plausibility. This aligns with recent research into "Geometric Hawkes Processes," which explicitly incorporate graph-based or other geometric structures to constrain the influence between processes, confirming the value of this synthesis.

Synthesis of Novelty and Comparative Table

The primary novelty of the Event-Time Geometry framework lies not in the invention of a single new mathematical object, but in its unique **synthesis** of concepts from disparate fields. It combines: (1) the continuous manifold and Minkowski-like metric from classical relativistic physics, providing a rigorous language for causality; (2) a probabilistic treatment of event coordinates, grounding the model in the reality of uncertain, noisy computational systems; and (3) a framework for imposing physical causal constraints on otherwise purely statistical models like STPPs.

ETG thus occupies a unique theoretical niche between purely logical models (Lamport), conservative bounded-uncertainty models (TrueTime), discrete fundamental physical models (CST), and purely statistical generative models (STPPs). The following table crystallizes these distinctions.

Feature	Lamport Clocks	Google TrueTime	Causal Set Theory	Event-Time Geometry (P-ETG)
Time Model	Logical (Counter)	Metric (Bounded Interval)	Acausal (Order Relation)	Metric (Probabilistic Distribution)
Structure	Topological (Partial	Continuous	Discrete (Poset)	Continuous

Feature	Lamport Clocks	Google TrueTime	Causal Set Theory	Event-Time Geometry (P-ETG)
	Order)	Manifold		Manifold
Causality	Binary (→)	Deterministic (Interval Non-overlap)	Binary (≺)	Probabilistic (P(e1 → e2))
Concurrency	Unordered	Resolved by waiting	Acausal Links	Probabilistically Ordered
Invariance	N/A	N/A	Lorentz Invariance (emergent)	Stochastic Lorentz-like Invariance
Primary Domain	Distributed Systems Theory	Distributed Databases	Quantum Gravity	Asynchronous Systems (Distributed, Neuromorphic)

This comparative analysis demonstrates that ETG provides a novel and valuable level of abstraction. It offers a formal, metric-aware, and uncertainty-conscious framework for reasoning about the very systems that are coming to dominate modern computing, providing a "middle way" that is both theoretically rich and pragmatically applicable.

Gap Analysis and Future Research Directions

While the Event-Time Geometry framework provides a powerful conceptual foundation, its current formulation is nascent and reveals several significant gaps and limitations. These unaddressed questions represent fertile ground for future research, pointing toward the extensions necessary to transform ETG from an elegant analogy into a comprehensive and robust theory of asynchronous computation. The framework is currently in a "classical" phase, analogous to physics before the development of field theories; the most profound future work involves treating the properties of the event-time space not as a fixed background, but as dynamic fields that evolve with the system's activity.

The Limits of the Special Relativistic Analogy

The current ETG framework is analogous only to Special Relativity, which describes physics in flat spacetime with inertial reference frames. This simplification, while useful for establishing the core concepts, limits its applicability to more complex, realistic systems.

Absence of "Curvature" (General Relativity)

D-ETG and P-ETG assume a "flat" event-time manifold, characterized by a globally constant maximum signal speed c and a static distance function d. However, real distributed systems exhibit highly dynamic and non-uniform performance characteristics. Network links can become congested, processors can become overloaded, and data centers can experience partitions. These phenomena effectively alter the "distance" between locations, slowing down information propagation.

This suggests a major research direction: developing a **"General Relativistic ETG"**. In such a framework, the event-time metric itself would be dynamic and non-uniform. Regions of high network latency or heavy computational load could be modeled as areas of high "curvature" in the event-time manifold, where the effective distance between points is greater and causal cones are consequently narrowed. This aligns with physics-inspired models of complex networks, such as network cosmology, which use geometric concepts to describe the large-scale structure and dynamics of systems like the Internet. A curved ETG would allow for a much more realistic model of system performance, where the geometry of causality is shaped by the flow of information and computation itself.

Meaning of "Velocity" and "Frames"

The Lorentz-like transformations in D-ETG introduce the concept of a relative "velocity" v between different event-time frames. This is used as a loose analogy for phenomena like clock skew or differences in processing speeds between two subsystems. However, this analogy lacks a rigorous, computationally grounded definition. What does it mean, precisely, for one clock domain to be "moving" relative to another? A key research gap is to formalize this concept by connecting it to measurable properties of computational systems. For example, relative velocity could be formally defined in terms of the first and second derivatives of the clock offset between two nodes, corresponding to clock drift and skew, respectively. Without such a grounding, the Lorentz-like transformations remain a compelling but incomplete metaphor.

The Nature of c

The maximum signal speed c is presented as a global constant, analogous to the speed of light in a vacuum. In modern heterogeneous computing systems, this assumption is unrealistic. The maximum speed of information propagation can vary by orders of magnitude between different components: on-chip communication via a shared bus is vastly faster than cross-datacenter communication over a wide-area network. The framework must be extended to handle a non-constant or location-dependent c(l_1, l_2), where the maximum speed is a function of the path between two locations. This would require moving beyond the simple Minkowski-like interval to a more general pseudo-Riemannian metric, further motivating the development of a General Relativistic ETG.

Deepening the Mathematical Connections

The P-ETG framework introduces connections to several powerful mathematical fields, but these connections are currently more conceptual than operational. Significant research is needed to transform these analogies into practical tools for analysis and design.

From Concept to Tool in Information Geometry

P-ETG correctly identifies Information Geometry as the natural language for measuring distances between probabilistic events ("event clouds"). However, this connection needs to be operationalized. A clear research direction is to use metrics from information geometry to formally define and quantify critical system properties. For instance, the **robustness of a temporal code** in a Spiking Neural Network (SNN) could be formally defined using the Fisher-Rao distance. A temporal code, which represents information via specific spike timings, is robust if the Fisher-Rao distance between the probability distributions of different spike patterns remains large even in the presence of spike-time jitter. This provides a quantitative, principled way to measure the information-carrying capacity of a noisy neural code. This approach would connect ETG directly to advanced theories of neural learning, such as natural-gradient-based plasticity, which are also formulated on a Riemannian manifold of probability distributions.

Modeling Uncertainty Accumulation with SDEs

P-ETG primarily describes the static uncertainty of an event's coordinates at a single point in time. A major gap is the lack of a model for *dynamic* uncertainty that evolves over time. Clock drift is a canonical example: the error in a local clock relative to a reference standard is not a fixed random variable but an accumulating process.
This can be rigorously modeled using the tools of **Stochastic Differential Geometry**. A concrete research program could involve defining the clock offset $\theta_i(t)$ of each node i as the solution to a stochastic differential equation (SDE), such as an Ornstein-Uhlenbeck process which models a random walk with a tendency to revert to a mean. For example:
where dW_t is a Wiener process term. By solving such SDEs, one can derive the probability distribution of the clock offset at any future time. This would allow for a principled analysis of how uncertainty accumulates along long causal chains, enabling the calculation of probabilities like, "What is the chance that this sequence of causally dependent tasks will miss its cumulative deadline due to accumulating jitter?".

Resolving Intransitivity in Probabilistic Ordering

The "likely-happened-before" relation, a natural consequence of P-ETG's probabilistic causal relation, is noted to be potentially intransitive. This is a significant theoretical weakness, as it means that pairwise comparisons can lead to logical cycles (e.g., $A \rightarrow B$, $B \rightarrow C$, but $C \rightarrow A$), preventing the construction of a consistent global partial order.
A key research gap is to develop a method for resolving this intransitivity. The connection to **Probabilistic Graphical Models (PGMs)** offers a promising path forward. Instead of relying on a series of independent pairwise comparisons, one could model the entire set of concurrent events as a single PGM. In this graph, each event's timestamp would be a random variable, and the causal probabilities would define the relationships between them. Standard inference algorithms, such as Markov Chain Monte Carlo (MCMC) sampling, could then be used to find the *globally most probable total ordering* of the events, conditioned on all available timing evidence. This approach would inherently resolve cycles and enforce transitivity, yielding a consistent, statistically optimal ordering where pairwise methods fail.

Unaddressed Operational and Algorithmic Questions

Beyond the theoretical foundations, the ETG framework lacks formal treatment of several key operational aspects required for its use in real-world system design and analysis.

Compositionality

The current framework describes the geometry of a single, monolithic system. A critical missing piece is a theory of **compositionality**. How do we formally compose two systems, each described by its own ETG? What is the resulting event-time manifold, and how are the causal structures combined? For example, if a neuromorphic vision system (System A) provides input to a distributed processing cluster (System B), a formal composition rule is needed to define the geometry of the combined system. Developing such a theory is essential for enabling modular design, verification, and analysis of large-scale, heterogeneous systems.

Practical Inference Algorithms

While P-ETG provides a formal definition of the causal influence probability, $P(e_1 \rightarrow e_2)$, as an integral over the events' joint distribution, it does not specify efficient algorithms for its computation. For simple cases like independent Gaussian distributions, this can be solved analytically. However, for more realistic scenarios involving non-Gaussian, correlated, or empirically derived distributions, the computation becomes intractable. Research is needed to develop practical and scalable algorithms for causal probability inference, potentially leveraging techniques from numerical integration, Monte Carlo methods (like importance sampling for rare events), or variational inference.

Learning from Data

The ETG framework is presented primarily as a modeling tool, where the parameters (e.g., jitter distributions, the effective value of c) are assumed to be known. A significant and valuable extension would be to develop algorithms to **learn the parameters of a P-ETG model from observational data**. Given a stream of timestamped events from a system, such as network packet logs or neural spike trains, techniques from statistical machine learning could be used to infer the underlying geometry. This would connect ETG to the rich field of causal inference from time series data, where methods are developed to uncover causal relationships from observational data, often within the framework of structural causal models or Granger causality. Learning the geometry from data would transform ETG from a purely analytical tool into an empirical one, allowing it to characterize and monitor the performance and causal structure of real, running systems.

A Guide to Formal Specification and Hardening

To mature from a compelling physical analogy into a rigorous theoretical framework suitable for verifying computational systems, Event-Time Geometry requires a process of "hardening." This involves developing a minimal, complete formal specification grounded in mathematical axioms and implemented within a proof assistant. This process is not merely a mechanical translation of ideas into formal language; it is a powerful tool for discovery that forces precision, exposes ambiguities in informal descriptions, and ultimately leads to a more robust and reliable theory. This section outlines a concrete four-step path toward this goal, leveraging established practices in formal methods for distributed systems.

Step 1: An Axiomatic Foundation for D-ETG

The first step is to move beyond the descriptive, analogy-driven presentation of D-ETG and define it as an abstract mathematical structure. This provides the minimal, complete foundation requested by the user. This can be achieved by defining an **Event-Time Space** as a tuple, for instance, $\mathcal{M} = (\mathcal{E}, L, T, d, c, \rightarrow)$. The properties of this space are then constrained by a set of axioms.

Axioms for Components

The components of the tuple are defined as sets with specific properties, which can be formally stated as axioms:

1. **Locations:** L is a set. The function $d: L \times L \rightarrow \mathbb{R}_{\ge 0}$ is a metric on L, satisfying the following axioms for all $l_1, l_2, l_3 \in L$:
   - Non-negativity: $d(l_1, l_2) \ge 0$
   - Identity of indiscernibles: $d(l_1, l_2) = 0 \iff l_1 = l_2$
   - Symmetry: $d(l_1, l_2) = d(l_2, l_1)$
   - Triangle inequality: $d(l_1, l_3) \le d(l_1, l_2) + d(l_2, l_3)$
2. **Timestamps:** T is a totally ordered set, equipped with an order relation $\le$, and is isomorphic to the real numbers $(\mathbb{R}, \le)$.
3. **Events:** $\mathcal{E}$ is a set. There exist projection functions $\pi_L: \mathcal{E} \rightarrow L$ and $\pi_T: \mathcal{E} \rightarrow T$ that map each event to its location and timestamp, respectively.

Axiom of Causality

The causal relation $\rightarrow \subseteq \mathcal{E} \times \mathcal{E}$ is not an independent component but is defined axiomatically based on the other components. For any two events $e_1, e_2 \in \mathcal{E}$:
where $c \in \mathbb{R}_{>0}$ is a constant of the space. This axiom elevates the causal inequality from a descriptive property to a definitional one.

Axiom of the Interval

The event-time interval function $s^2: \mathcal{E} \times \mathcal{E} \rightarrow \mathbb{R}$ is defined as:
With this axiomatic foundation, the Event-Time Invariance Theorem is no longer a statement to be taken on faith from the physics analogy, but a formal theorem to be derived. The proof would involve formally defining the class of Lorentz-like transformations on the space and demonstrating that the function s^2 is invariant under their application.

Step 2: Formalizing the Probabilistic Extension (P-ETG)

Building upon the axiomatic foundation of D-ETG, the probabilistic extension can be formalized by leveraging the mathematical theory of probability spaces and random variables.

Defining the Probabilistic Event Space

The concept of an event must be elevated from a point to a formal probabilistic object. An **Event-Time Manifold** is first defined as the measurable space $(L \times T, \mathcal{B})$, where $\mathcal{B}$ is a suitable σ-algebra (e.g., the Borel σ-algebra if L and T have topological structure). A **probabilistic event** E is then formally defined as a probability measure on this space. The set of all possible events is thus a subset of the space of all such probability measures.

Formal Derivation of Stochastic Quantities

With events defined as probability measures, the derived quantities become formally defined random variables. Let E_1 and E_2 be two independent probabilistic events with corresponding random variables (L_1, T_1) and (L_2, T_2).

- The **temporal separation** ΔT is the random variable $T_2 - T_1$. Its distribution is the convolution of the distribution of T_2 and the distribution of $-T_1$.
- The **spatial separation** ΔX is the random variable $d(L_1, L_2)$. Its distribution must be derived from the joint distribution of L_1 and L_2.
- The **stochastic interval** $(\Delta S)^2$ is the random variable $c^2(\Delta T)^2 - (\Delta X)^2$. Its probability distribution can, in principle, be derived using the calculus of random variables. For example, if the event coordinates are independent and follow Gaussian distributions, the distribution of $(\Delta S)^2$ would be related to a non-central chi-squared distribution, which has a precise mathematical form.

Formalizing the Probabilistic Causal Relation

The causal influence probability $P(E_1 \rightarrow E_2)$ is formally defined as an integral over the joint probability distribution of the two events. Let $P_{1,2}(l_1, t_1, l_2, t_2)$ be the joint probability density function for the coordinates of E_1 and E_2. Then:
where the region of integration R is the set of all coordinate tuples (l_1, t_1, l_2, t_2) that satisfy the causal inequality $t_2 \ge t_1 + d(l_1, l_2)/c$. The properties of this relation, such as its potential for intransitivity, can then be formally investigated and proven.

Step 3: A Roadmap for Implementation in a Proof Assistant (Coq)

The formal axiomatic system provides a direct blueprint for implementation and verification in a proof assistant like Coq, which is well-suited for specifying and verifying properties of distributed systems.

Representing the Axioms

The axiomatic system from Step 1 can be directly encoded using Coq's powerful type theory.

- The set of locations L can be defined as a Coq Type. The metric d can be defined as a function on this type, with the metric axioms stated and proven as lemmas about this function.
- Events can be represented as a record type (a struct) containing fields for a location and a timestamp.
- The causal relation $\rightarrow$ and the interval function s^2 can be implemented as Coq functions or propositions that directly mirror their axiomatic definitions.

Proving the Invariance Theorem

The Event-Time Invariance Theorem would become a key goal for formal proof within the system. This process would involve:

1. Defining the set of Lorentz-like transformations as Coq functions.

2. Stating the theorem: forall (T: Transformation) (e1 e2: Event), is_lorentz_like(T) -> interval(e1, e2) = interval(T(e1), T(e2)).
3. Interactively constructing the proof in Coq.

Successfully completing this proof would provide an extremely high-confidence guarantee of the internal consistency of the D-ETG framework. The process of formalization itself would act as a discovery tool. For instance, to prove the theorem, one must provide a perfectly precise definition of what constitutes a valid "inertial frame" and a "Lorentz-like transformation" in a computational context. This forces the replacement of the loose physical analogy with a rigorous mathematical definition, thereby hardening the theory by exposing and resolving ambiguities.

Developing a P-ETG Library

Formalizing P-ETG in Coq is a more advanced task that would require leveraging a library for probability theory or measure theory.

- The "event clouds" would be represented not as simple records, but as formal probability distributions or measures.
- The key deliverable of this effort would be a verified function, for example, compute_causal_prob(E1, E2). This function would take two probabilistic events (represented as probability measures) as input and return a real number representing their causal influence probability.
- The correctness theorem for this function would state that the value it computes is provably equal to the formal integral definition from Step 2.

Frameworks like Verdi, which is used to verify distributed system implementations in Coq, provide a model for how such a verification effort could be structured.

Step 4: From Formal Model to Verifiable Protocols

The ultimate purpose of developing a formal specification is to enable the verification of real systems and protocols that operate under the assumptions of the model. The hardened ETG framework can serve as the semantic foundation for specifying and verifying the correctness of complex, time-sensitive protocols.

Example: Verifying a Probabilistic Fair Ordering Protocol

Consider a probabilistic fair ordering protocol, such as the "Tommy" system, which aims to order concurrent events based on the "likely-happened-before" relation. A formal verification of such a protocol using the P-ETG framework would proceed as follows:

1. **Specify the Protocol:** The state of each node and the network, as well as the rules for handling messages and updating state, would be formally specified as an automaton or state machine within the Coq framework, similar to approaches used in CoqIOA.
2. **Define the Fairness Property:** The desired "fairness" property would be formally stated using the P-ETG causal relation. For example: "For any two events E_1, E_2 generated by the clients, if the P-ETG model of the system implies that P(E_1 \rightarrow E_2) > 1 - \epsilon, then the protocol must ensure that its output sequence orders e_1 before e_2 with a probability greater than 1 - \delta."
3. **Prove Correctness:** The final, and most challenging, step would be to construct a formal proof in Coq that the protocol specification satisfies the fairness property. This would involve reasoning about all possible executions of the protocol under the probabilistic

assumptions of the P-ETG model.

Successfully completing such a proof would provide an unprecedented level of assurance for the fairness and correctness of such protocols, moving them from heuristic designs to verifiably robust systems.

Publication Strategy and Journal Targeting

A novel and multifaceted framework like Event-Time Geometry requires a staged and strategic approach to publication to maximize its impact across relevant scientific communities. A single, monolithic paper attempting to cover its theoretical depth and practical breadth would likely fail, as it would be too theoretical for systems researchers and too applied for pure theorists. A successful research program should therefore be structured as a portfolio of publications, each tailored to a specific audience and contribution type. This section provides a guide to targeting appropriate journals and conferences for disseminating the various facets of ETG research.

Foundational Theory and Formal Methods (Targeting the Core of TCS)

This stream of work focuses on the fundamental mathematical and logical underpinnings of the ETG framework. The contributions are abstract, rigorous, and aimed at the theoretical computer science (TCS) community.

- **Contribution Type:** This includes papers presenting the axiomatic foundations of D-ETG and P-ETG (as outlined in Section 4.1 and 4.2), the formal, machine-checked proof of the Event-Time Invariance Theorem, and the resolution of deep theoretical gaps, such as the intransitivity of the probabilistic causal relation (as discussed in Section 3.2).
- **Target Venues:** The most appropriate venues are those that value deep theoretical contributions, formal rigor, and foundational models of computation.
    - **Top-Tier Journals:** The premier journals for this type of work are the *Journal of the ACM* , the *SIAM Journal on Computing* , *Theoretical Computer Science* , and *Information and Computation*. These journals have a long history of publishing seminal work in the theory of computation and distributed systems.
    - **Premier Conferences:** The most prestigious conferences for presenting foundational theoretical results are the *ACM Symposium on Theory of Computing (STOC)* , the *IEEE Symposium on Foundations of Computer Science (FOCS)* , and for work with a strong logical or formal methods component, the *IEEE Symposium on Logic in Computer Science (LICS)*. Acceptance at these venues would firmly establish the theoretical credibility of the ETG framework.

Systems Applications (Targeting Domain Specialists)

This stream of work focuses on applying the P-ETG framework to solve concrete, pressing problems in specific application domains, demonstrating the practical utility of the theory. The narrative of these papers should be driven by the problem and the solution, with the ETG framework presented as the key enabling tool.

For Distributed Systems

- **Contribution Type:** Research that designs, implements, and evaluates a P-ETG-based

protocol. Examples include a protocol for probabilistic fair ordering that outperforms centralized sequencers, a resilient scheduling algorithm that uses probabilistic deadlines to improve resource utilization, or an anomaly detection system that uses "Causal Fidelity" to monitor system health.
- **Target Venues:** The target audience consists of researchers and practitioners in distributed systems, operating systems, and networking.
  - **Top Journals:** *IEEE Transactions on Parallel and Distributed Systems* is a flagship journal in the field and an ideal target. Other strong candidates include the *Journal of Parallel and Distributed Computing* and *Distributed Computing*.
  - **Top Conferences:** The premier venues for impactful systems research are the *USENIX Symposium on Operating Systems Design and Implementation (OSDI)* and the *USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. Presenting at these conferences would demonstrate the practical relevance and performance benefits of ETG-based systems.

For Neuromorphic Computing

- **Contribution Type:** Research that uses P-ETG to model and analyze neuromorphic systems. Examples include a quantitative analysis of the robustness of temporal codes to spike-time jitter, a novel algorithm for causal inference in noisy spike trains that leverages the causal cone constraint, or a new uncertainty-aware spike-timing-dependent plasticity (STDP) rule inspired by the P-ETG formalism.
- **Target Venues:** This work should be directed at the specialized community of researchers in neuromorphic engineering, computational neuroscience, and brain-inspired AI.
  - **Specialized Journals:** The most direct path to this community is through dedicated journals such as *Neuromorphic Computing and Engineering* and the *Neuromorphic Engineering* section of *Frontiers in Neuroscience*. These journals are edited and reviewed by experts in the field and ensure the work reaches the intended audience.

High-Impact Interdisciplinary Work (Bridging Physics, CS, and AI)

This stream represents the most ambitious and potentially highest-impact publications. These are seminal papers that either introduce a major theoretical advance with broad implications or demonstrate a breakthrough application that unifies theory and practice in a compelling and accessible way.

- **Contribution Type:** A paper introducing a "General Relativistic ETG" that models dynamic latency as spacetime curvature would fall into this category. Similarly, a paper that not only proposes a P-ETG-based fair ordering protocol but also demonstrates its successful deployment in a real-world, high-stakes application (like a financial exchange) could have broad appeal. The key is to frame the work not just as an incremental advance in a subfield, but as a new way of thinking about a fundamental problem.
- **Target Venues:** The goal is to reach a wide scientific audience beyond a single discipline.
  - **Broad Scientific Journals:** High-impact, multidisciplinary journals such as *Proceedings of the National Academy of Sciences (PNAS)* and *Nature Communications* are appropriate targets. A paper successfully framing ETG as a new "physics of computation" that unifies ideas from relativity, probability theory,

and computer science would be a strong candidate.
   - **Interdisciplinary Physics/CS Journals:** Journals that explicitly welcome cross-disciplinary work, such as *Chaos: An Interdisciplinary Journal of Nonlinear Science* or *Insights in Theoretical, Computational, and Interdisciplinary Physics* , provide another avenue for reaching a mixed audience of physicists and computer scientists.

This staged, multi-venue strategy allows the research program to build momentum. The foundational theoretical publications provide a solid, citable basis for the framework. The domain-specific application papers then demonstrate its value and utility to expert communities. Finally, the high-impact interdisciplinary papers can synthesize these achievements to introduce the ETG framework to the broader scientific world as a mature and powerful new paradigm.

Conclusion

This analysis has charted the evolution of Event-Time Geometry from its deterministic origins, rooted in an analogy to relativistic spacetime, to a comprehensive probabilistic framework capable of modeling the inherent uncertainty of real-world computational systems. The journey from D-ETG to P-ETG is marked by a systematic transformation: precise event coordinates become probability distributions or "event clouds"; sharp causal boundaries dissolve into fuzzy, probabilistic gradients; and absolute causal relations are replaced by graded measures of confidence. This extension is not merely an addition of noise to an existing model but a fundamental reconceptualization that yields a richer, more physically grounded, and ultimately more applicable theory of causality in complex systems.

The novelty of ETG is amplified by its unique position at the intersection of several fields. It advances beyond the purely logical ordering of Lamport clocks by introducing a metric, and it offers a more nuanced view of uncertainty than the conservative, worst-case bounds of systems like TrueTime. By modeling the full distribution of uncertainty, it enables a new class of algorithms and protocols that can make statistically optimal decisions, balancing performance and reliability based on quantifiable risk. Furthermore, its connection to established mathematical disciplines—including Information Geometry for measuring the distance between uncertain events, Stochastic Differential Geometry for modeling the dynamics of accumulating uncertainty, and Probabilistic Graphical Models for performing complex inference—provides a robust theoretical underpinning that ensures P-ETG is a principled framework unifying geometry, probability, and information theory.

The practical benefits of this framework are most evident in two of the most challenging domains of modern computer science. In neuromorphic systems, where computation is encoded in the precise timing of stochastic spike events, P-ETG provides the essential tools to analyze the robustness of temporal codes, infer causal connectivity from noisy data, and design systems that are resilient to the inherent jitter of their physical substrate. In distributed systems, P-ETG offers a foundation for solving long-standing problems in fairness, performance, and resilience. It enables the design of probabilistically fair ordering protocols that eliminate the arbitrariness of centralized sequencers, facilitates high-performance probabilistic scheduling that avoids the over-provisioning of deterministic methods, and provides novel metrics like Causal Fidelity for monitoring system integrity.

Looking forward, the identified research gaps—such as developing a "General Relativistic" extension to model dynamic latencies, operationalizing the connections to information geometry, and creating algorithms to learn the geometry from data—chart a clear path for future work. The process of formalizing the framework in a proof assistant will be critical to this endeavor, serving not only to guarantee its internal consistency but also to refine its core concepts by replacing loose analogy with mathematical precision. As computational systems become increasingly decentralized, asynchronous, and intertwined with the stochastic physical world, the ability to reason rigorously under temporal and spatial uncertainty will become a central design challenge. Probabilistic geometries like P-ETG represent a necessary theoretical foundation for this next generation of computing, extending the elegant structure of spacetime to the domain of stochastic information processing and ensuring that even as we embrace randomness, we retain a consistent and quantifiable geometry of cause and effect.

Works cited

1. Beyond Lamport, Towards Probabilistic Fair Ordering - arXiv, https://www.arxiv.org/pdf/2510.13664 2. The Mutual Exclusion Problem Part I: A Theory of Interprocess Communication - Leslie Lamport, https://lamport.azurewebsites.net/pubs/mutual1.pdf 3. Fair Ordering - arXiv, https://arxiv.org/pdf/2510.13664 4. Virtual Time and Global States of Distributed Systems, https://vs.inf.ethz.ch/publ/papers/VirtTimeGlobStates.pdf 5. Virtual Time and Global States of Distributed Systems - csail, https://courses.csail.mit.edu/6.852/01/papers/VirtTime_GlobState.pdf 6. My Takeaways from the Original Google Spanner Paper | by Akshat Jaimini | Medium, https://medium.com/@destrex271/my-takeaways-from-the-original-google-spanner-paper-a3b2683952fe 7. Beyond TrueTime: Using AugmentedTime for Improving Spanner, https://cse.buffalo.edu/~demirbas/publications/augmentedTime.pdf 8. Notes on the Google Spanner Paper | Lu's blog, https://blog.the-pans.com/notes-on-the-spanner/ 9. arXiv:1203.2109v2 [gr-qc] 26 Nov 2012 - CAIDA.org, https://www.caida.org/catalog/papers/2012_network_cosmology/network_cosmology.pdf 10. [2410.02862] Do causal sets have symmetries? - arXiv, https://arxiv.org/abs/2410.02862 11. Retarded Causal Set Propagator in 2D Anti de-Sitter Spacetime, https://arxiv.org/pdf/2504.12919 12. Causal sets - Wikipedia, https://en.wikipedia.org/wiki/Causal_sets 13. Probabilistic Spacetimes, https://arxiv.org/pdf/1712.06127 14. Neural Temporal Point Processes: A Review - IJCAI, https://www.ijcai.org/proceedings/2021/0623.pdf 15. Deep Spatio-temporal Point Processes: Advances and New Directions 1 - arXiv, https://arxiv.org/html/2504.06364v1 16. Modeling Continuous-time Event Data with [2mm] Neural Temporal Point Processes - mediaTUM, https://mediatum.ub.tum.de/doc/1662914/2wdsxe2av36cxz519qo7xolo3.dissertation.pdf 17. Hawkes Models and Their Applications - Annual Reviews, https://www.annualreviews.org/content/journals/10.1146/annurev-statistics-112723-034304 18. Causal Discovery in Hawkes Processes by Minimum Description Length, https://ojs.aaai.org/index.php/AAAI/article/view/20656/20415 19. A Brief Introduction to Hawkes Processes | by Hey Amit | We Talk Data | Medium, https://medium.com/we-talk-data/a-brief-introduction-to-hawkes-processes-996d8da8e003 20. Machine learning based on Hawkes processes and stochastic optimization - ResearchGate, https://www.researchgate.net/publication/336601787_Machine_learning_based_on_Hawkes_processes_and_stochastic_optimization 21. Natural-gradient learning for spiking neurons - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC9038192/ 22. CITS: Nonparametric Statistical Causal Modeling for High-Resolution Neural Time Series, https://arxiv.org/html/2508.01920v1 23. Estimating the directed information to infer causal relationships in ensemble neural spike train

recordings - NIH, https://pmc.ncbi.nlm.nih.gov/articles/PMC3171872/ 24. uwplse/verdi: A framework for formally verifying distributed systems implementations in Coq - GitHub, https://github.com/uwplse/verdi 25. CoqIOA: A Formalization of IO Automata in the Coq Proof Assistant - acm sigops, https://www.sigops.org/src/srcsosp2017/sosp17src-final41.pdf 26. Formalization of Quantum Protocols using Coq - arXiv, https://arxiv.org/pdf/1511.01568 27. CaitieM20/TheVerificationOfDistributedSystem: Accompanying Repository for The Verification of a Distributed System Talk - GitHub, https://github.com/CaitieM20/TheVerificationOfDistributedSystem 28. Theoretical Computer Science - Google Scholar Metrics, https://scholar.google.com/citations?view_op=top_venues&hl=en&vq=eng_theoreticalcomputerscience 29. Computer Science, Theory & Methods: Journal Rankings | OOIR, https://ooir.org/journals.php?field=Computer+Science&category=Computer+Science%2C+Theory+%26+Methods 30. CORE Computer Science Journal Rankings, https://cic.tju.edu.cn/faculty/zhileiliu/doc/COREComputerScienceJournalRankings.html 31. Computing Systems - Google Scholar Metrics, https://scholar.google.com/citations?view_op=top_venues&hl=en&vq=eng_computingsystems 32. IEEE Transactions on Parallel and Distributed Systems | IEEE Computer Society Digital Library, https://www.computer.org/csdl/journal/td 33. www.google.com, https://www.google.com/search?q=top+journals+distributed+systems 34. Top 18 Distributed Computing papers published in 2023 - SciSpace, https://scispace.com/journals/distributed-computing-2lho336i/2023 35. Journals - - NeuroPAC, https://www.neuropac.info/journals/ 36. About Neuromorphic Computing and Engineering - IOPscience - Publishing Support, https://publishingsupport.iopscience.iop.org/journals/neuromorphic-computing-engineering/about-neuromorphic-computing-engineering/ 37. Neuromorphic Computing and Engineering - Impact Factor & Score 2025 | Research.com, https://research.com/journal/neuromorphic-computing-and-engineering 38. Frontiers in Neuroscience | Neuromorphic Engineering, https://www.frontiersin.org/journals/neuroscience/sections/neuromorphic-engineering 39. PNAS – Explore High-Impact Scientific Research Across Disciplines from One of the World's Most-Cited Journals, https://www.pnas.org/ 40. Journals - AIP Publishing, https://pubs.aip.org/pages/journals 41. Insights in Theoretical, Computational, and Interdisciplinary Physics - Lanash Science, https://lanashscience.com/journal/insights-in-theoretical-computational-and-interdisciplinary-physics